

Computer Networks Lab

Socket Programming

Prof. Nikita Singhal
Department of Computer Engineering

Contents

2

- Client-Server Communication
- Socket and Socket based communication
- Socket Characteristics
- Socket API
- Connection based communication
- Connection less communication
- Socket Programming

Client Server Communication

3

□ Server

- passively waits for and responds to clients
- passive socket

□ Client

- initiates the communication
- must know the address and the port of the server
- active socket

Sockets and Socket-based Communication

4

- ❑ Sockets provide an interface for programming networks at the transport layer. Network communication using Sockets is very much similar to performing file I/O.
- ❑ The streams used in file I/O operation are also applicable to socket-based I/O.
- ❑ Socket-based communication is independent of a programming language used for implementing it.
- ❑ That means, a socket program written in Java language can communicate to a program written in non-Java (say C or C++) socket program.

Socket Characteristics

5

- Socket are characterized by their domain, type and transport protocol.
- Common domains are:
 - ▣ IAF UNIX: address format is UNIX pathname
 - ▣ IAF INET: address format is host and port number
- Common types are:
 - ▣ virtual circuit: received in order transmitted and reliably
 - ▣ datagram: arbitrary order, unreliable

Socket Characteristics(continued)

6

- Each socket type has one or more protocols. Ex:
 - ▣ TCP/IP (virtual circuits)
 - ▣ UDP (datagram)
- Use of sockets:
 - ▣ Connection-based sockets communicate client-server: the server waits for a connection from the client
 - ▣ Connectionless sockets are peer-to-peer: each process is symmetric.

Socket Family

7

Name	Purpose
AF_UNIX, AF_LOCAL	Local communication
AF_INET	IPv4 Internet protocols
AF_INET6	IPv6 Internet protocols
AF_IPX	IPX - Novell protocols
AF_NETLINK	Kernel user interface device
AF_X25	ITU-T X.25 / ISO-8208 protocol
AF_AX25	Amateur radio AX.25 protocol
AF_ATMPVC	Access to raw ATM PVCs
AF_APPLETALK	Appletalk
AF_PACKET	Low level packet interface

Socket APIs

8

- `socket`: creates a socket of a given domain, type, protocol (buy a phone)
- `bind`: assigns a name to the socket (get a telephone number)
- `listen`: specifies the number of pending connections that can be queued for a server socket. (call waiting allowance)
- `accept`: server accepts a connection request from a client (answer phone)
- `connect`: client requests a connection request to a server (call)
- `send`, `sendto`: write to connection (speak)
- `recv`, `recvfrom`: read from connection (listen)
- `shutdown`: end the call

Connection-based communication (TCP)

9

- Server performs the following actions
 - ▣ socket: create the socket
 - ▣ bind: give the address of the socket on the server
 - ▣ listen: specifies the maximum number of connection requests that can be pending for this process
 - ▣ accept: establish the connection with a specific client
 - ▣ send,recv: stream-based equivalents of read and write (repeated)
 - ▣ shutdown: end reading or writing
 - ▣ close: release kernel data structures

Connection-based communication (TCP)

10

- Client performs the following actions
 - socket: create the socket
 - connect: connect to a server
 - send, recv: (repeated)
 - shutdown
 - close

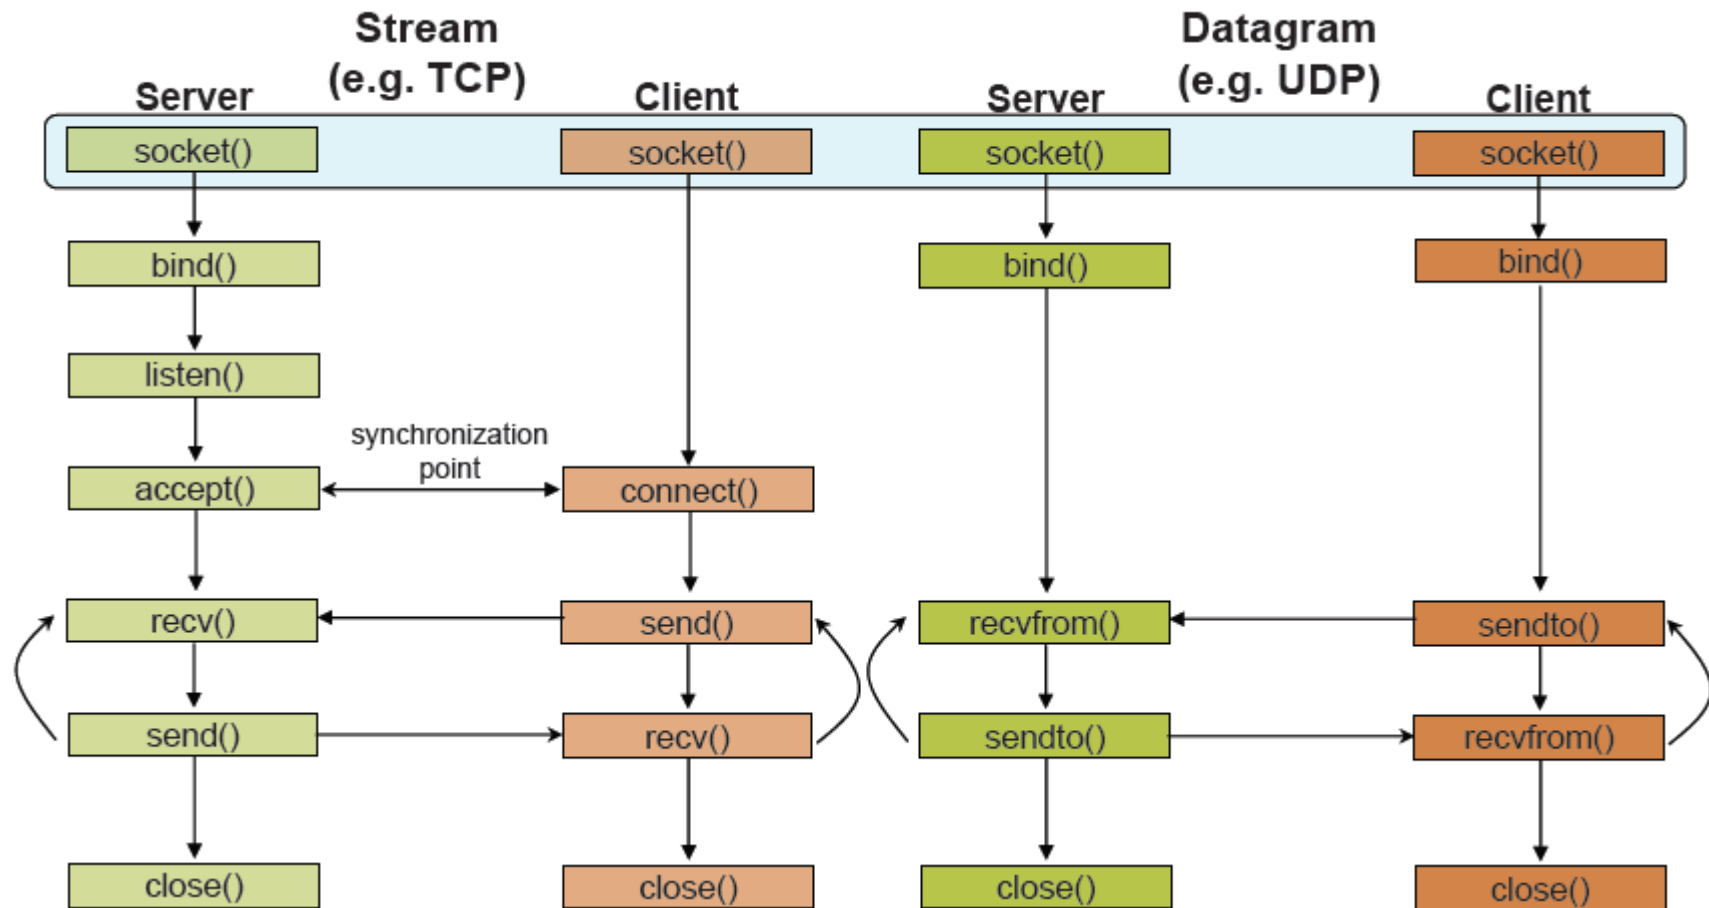
Connection-less communication (UDP)

11

- Communication is symmetric (peer-to-peer)
 - socket
 - bind: bind is optional for initiator
 - sendto, recvfrom (repeated)
 - shutdown
 - close

Client-Server Communication (Socket Programming)

12



Socket creation in C: socket()

- `int sockid = socket(family, type, protocol);`
 - **sockid**: socket descriptor, an integer (like a file-handle)
 - **family**: integer, communication domain, e.g.,
 - PF_INET, IPv4 protocols, Internet addresses (typically used)
 - PF_UNIX, Local communication, File addresses
 - **type**: communication type
 - SOCK_STREAM - reliable, 2-way, connection-based service
 - SOCK_DGRAM - unreliable, connectionless, messages of maximum length
 - **protocol**: specifies protocol
 - IPPROTO_TCP IPPROTO_UDP
 - usually set to 0 (i.e., use default protocol)
 - upon failure returns -1
- ☞ NOTE: socket call does not specify where data will be coming from, nor where it will be going to – it just creates the interface!

Specifying Addresses

14

- Socket API defines a **generic** data type for addresses:

```
struct sockaddr {  
    unsigned short sa_family; /* Address family (e.g. AF_INET) */  
    char sa_data[14];        /* Family-specific address information */  
}
```

- Particular form of the sockaddr used for **TCP/IP** addresses:

```
struct in_addr {  
    unsigned long s_addr; /* Internet address (32 bits) */  
}  
  
struct sockaddr_in {  
    unsigned short sin_family; /* Internet protocol (AF_INET) */  
    unsigned short sin_port; /* Address port (16 bits) */  
    struct in_addr sin_addr; /* Internet address (32 bits) */  
    char sin_zero[8]; /* Not used */  
}
```

👉 **Important:** sockaddr_in can be casted to a sockaddr

Assign address to socket: bind()

15

- associates and reserves a port for use by the socket
- `int status = bind(sockid, &addrport, size);`
 - **sockid**: integer, socket descriptor
 - **addrport**: struct sockaddr, the (IP) address and port of the machine
 - for TCP/IP server, internet address is usually set to INADDR_ANY, i.e., chooses any incoming interface
 - **size**: the size (in bytes) of the addrport structure
 - **status**: upon failure -1 is returned

bind()-Example with TCP

16

```
int sockid;
struct sockaddr_in addrport;
sockid = socket(PF_INET, SOCK_STREAM, 0);

addrport.sin_family = AF_INET;
addrport.sin_port = htons(5100);
addrport.sin_addr.s_addr = htonl(INADDR_ANY);
if(bind(sockid, (struct sockaddr *) &addrport, sizeof(addrport)) != -1) {
    ...}
```


listen()

17

- Instructs TCP protocol implementation to listen for connections
- ```
int status = listen(sockid, queueLimit);
```

  - **sockid**: integer, socket descriptor
  - **queueLen**: integer, # of active participants that can “wait” for a connection
  - **status**: 0 if listening, -1 if error
- `listen()` is **non-blocking**: returns immediately
- The listening socket (sockid)
  - is never used for sending and receiving
  - is used by the server only as a way to get new sockets

# Establish Connection: connect()

18

- The client establishes a connection with the server by calling `connect()`
- `int status = connect(sockid, &foreignAddr, addrlen);`
  - `sockid`: integer, socket to be used in connection
  - `foreignAddr`: struct `sockaddr`: address of the passive participant
  - `addrlen`: integer, `sizeof(name)`
  - status: 0 if successful connect, -1 otherwise
- `connect()` is **blocking**

# Incoming Connection: accept()

19

- The server gets a socket for an incoming client connection by calling `accept()`
- `int s = accept(sockid, &clientAddr, &addrLen);`
  - `s`: integer, the new socket (used for data-transfer)
  - `sockid`: integer, the orig. socket (being listened on)
  - `clientAddr`: struct `sockaddr`, address of the active participant
    - filled in upon return
  - `addrLen`: `sizeof(clientAddr)`: value/result parameter
    - must be set appropriately before call
    - adjusted upon return
- `accept()`
  - is **blocking**: waits for connection before returning
  - dequeues the next connection on the queue for socket (`sockid`)

# Exchanging data with stream socket(TCP)

20

- `int count = send(sockid, msg, msgLen, flags);`
  - `msg`: const void[], message to be transmitted
  - `msgLen`: integer, length of message (in bytes) to transmit
  - `flags`: integer, special options, usually just 0
  - `count`: # bytes transmitted (-1 if error)
- `int count = recv(sockid, recvBuf, bufLen, flags);`
  - `recvBuf`: void[], stores received bytes
  - `bufLen`: # bytes received
  - `flags`: integer, special options, usually just 0
  - `count`: # bytes received (-1 if error)
- Calls are **blocking**
  - returns only after data is sent / received

# Exchanging data with datagram socket (UDP)

21

- `int count = sendto(sockid, msg, msgLen, flags, &foreignAddr, addrlen);`
  - `msg, msgLen, flags`, count: same with `send()`
  - `foreignAddr`: struct `sockaddr`, address of the destination
  - `addrlen`: `sizeof(foreignAddr)`
- `int count = recvfrom(sockid, recvBuf, bufLen, flags, &clientAddr, addrlen);`
  - `recvBuf, bufLen, flags`, count: same with `recv()`
  - `clientAddr`: struct `sockaddr`, address of the client
  - `addrlen`: `sizeof(clientAddr)`
- Calls are **blocking**
  - returns only after data is sent / received

# Socket close in C: close()

22

- When finished using a socket, the socket should be closed
- `status = close(sockid);`
  - **sockid**: the file descriptor (socket being closed)
  - **status**: 0 if successful, -1 if error
- Closing a socket
  - closes a connection (for stream socket)
  - frees up the port used by the socket

# Example -Echo

23

- A client communicates with an “echo” server
- The server simply echoes whatever it receives back to the client

# Example -Echo using stream socket

24

The server starts by getting ready to receive client connections...

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. Accept new connection
  - b. Communicate
  - c. Close the connection



# Example -Echo using stream socket

25

```
/* Create socket for incoming connections */
if ((servSock = socket(PF_INET, SOCK_STREAM, IPPROTO_TCP)) < 0)
 DieWithError("socket() failed");
```

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. **Create a TCP socket**
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. Accept new connection
  - b. Communicate
  - c. Close the connection

# Example -Echo using stream socket

26

```
echoServAddr.sin_family = AF_INET; /* Internet address family */
echoServAddr.sin_addr.s_addr = htonl(INADDR_ANY); /* Any incoming interface */
echoServAddr.sin_port = htons(echoServPort); /* Local port */

if (bind(servSock, (struct sockaddr *) &echoServAddr, sizeof(echoServAddr)) < 0)
 DieWithError("bind() failed");
```

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. **Assign a port to socket**
3. Set socket to listen
4. Repeatedly:
  - a. Accept new connection
  - b. Communicate
  - c. Close the connection

# Example -Echo using stream socket

27

```
/* Mark the socket so it will listen for incoming connections */
if (listen(servSock, MAXPENDING) < 0)
 DieWithError("listen() failed");
```

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. **Set socket to listen**
4. Repeatedly:
  - a. Accept new connection
  - b. Communicate
  - c. Close the connection

# Example -Echo using stream socket

28

```
for (;;) /* Run forever */
{
 clntLen = sizeof(echoClntAddr);

 if ((clientSock=accept(servSock, (struct sockaddr *)&echoClntAddr, &clntLen))<0)
 DieWithError("accept() failed");
 ...
}
```

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. **Accept new connection**
  - b. Communicate
  - c. Close the connection

# Example -Echo using stream socket

29

Server is now blocked waiting for connection from a client

...

A client decides to talk to the server

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. **Accept new connection**
  - b. Communicate
  - c. Close the connection

# Example -Echo using stream socket

30

```
/* Create a reliable, stream socket using TCP */
if ((clientSock = socket(PF_INET, SOCK_STREAM, IPPROTO_TCP)) < 0)
 DieWithError("socket() failed");
```

## Client

1. **Create a TCP socket**
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. **Accept new connection**
  - b. Communicate
  - c. Close the connection

# Example -Echo using stream socket

31

```
echoServAddr.sin_family = AF_INET; /* Internet address family */
echoServAddr.sin_addr.s_addr = inet_addr(echoservIP); /* Server IP address*/
echoServAddr.sin_port = htons(echoServPort); /* Server port */

if (connect(clientSock, (struct sockaddr *) &echoServAddr,
 sizeof(echoServAddr)) < 0)
 DieWithError("connect() failed");
```

## Client

1. Create a TCP socket
2. **Establish connection**
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. **Accept new connection**
  - b. Communicate
  - c. Close the connection

# Example -Echo using stream socket

32

Server's accept procedure is now unblocked and returns client's socket

```
for (;;) /* Run forever */
{
 clntLen = sizeof(echoClntAddr);

 if ((clientSock=accept(servSock, (struct sockaddr *)&echoClntAddr, &clntLen))<0)
 DieWithError("accept() failed");
 ...
}
```

## Client

1. Create a TCP socket
2. **Establish connection**
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. **Accept new connection**
  - b. Communicate
  - c. Close the connection



# Example -Echo using stream socket

33

```
echoStringLen = strlen(echoString); /* Determine input length */

/* Send the string to the server */
if (send(clientSock, echoString, echoStringLen, 0) != echoStringLen)
 DieWithError("send() sent a different number of bytes than expected");
```

## Client

1. Create a TCP socket
2. Establish connection
3. **Communicate**
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. **Accept new connection**
  - b. Communicate
  - c. Close the connection

# Example -Echo using stream socket

34

```
/* Receive message from client */
if ((recvMsgSize = recv(clntSocket, echoBuffer, RCVBUFSIZE, 0)) < 0)
 DieWithError("recv() failed");
/* Send received string and receive again until end of transmission */
while (recvMsgSize > 0) { /* zero indicates end of transmission */
 if (send(clientSocket, echobuffer, recvMsgSize, 0) != recvMsgSize)
 DieWithError("send() failed");
 if ((recvMsgSize = recv(clientSocket, echoBuffer, RCVBUFSIZE, 0)) < 0)
 DieWithError("recv() failed");
}
```

## Client

1. Create a TCP socket
2. Establish connection
3. **Communicate**
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. Accept new connection
  - b. **Communicate**
  - c. Close the connection

# Example -Echo using stream socket

35

Similarly, the client receives the data from the server

## Client

1. Create a TCP socket
2. Establish connection
3. **Communicate**
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. Accept new connection
  - b. **Communicate**
  - c. Close the connection

# Example -Echo using stream socket

36

```
close(clientSock);
```

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. **Close the connection**

```
close(clientSock);
```

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. Accept new connection
  - b. Communicate
  - c. **Close the connection**

# Example -Echo using stream socket

37

Server is now blocked waiting for connection from a client

...

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. **Accept new connection**
  - b. Communicate
  - c. Close the connection

# Example -Echo using Datagram socket

38

```
/* Create socket for sending/receiving datagrams */
if ((servSock = socket(PF_INET, SOCK_DGRAM, IPPROTO_UDP)) < 0)
 DieWithError("socket() failed");
```

```
/* Create a datagram/UDP socket */
if ((clientSock = socket(PF_INET, SOCK_DGRAM, IPPROTO_UDP)) < 0)
 DieWithError("socket() failed");
```

## Client

1. **Create a UDP socket**
2. Assign a port to socket
3. Communicate
4. Close the socket

## Server

1. **Create a UDP socket**
2. Assign a port to socket
3. Repeatedly
  - Communicate

# Example -Echo using Datagram socket

39

```
echoServAddr.sin_family = AF_INET; /* Internet address family */
echoServAddr.sin_addr.s_addr = htonl(INADDR_ANY); /* Any incoming interface */
echoServAddr.sin_port = htons(echoServPort); /* Local port */

if (bind(servSock, (struct sockaddr *) &echoServAddr, sizeof(echoServAddr)) < 0)
 DieWithError("bind() failed");
```

```
echoClientAddr.sin_family = AF_INET; /* Internet address family */
echoClientAddr.sin_addr.s_addr = htonl(INADDR_ANY); /* Any incoming interface */
echoClientAddr.sin_port = htons(echoClientPort); /* Local port */

if (bind(clientSock, (struct sockaddr *) &echoClientAddr, sizeof(echoClientAddr)) < 0)
 DieWithError("connect() failed");
```

## Client

1. Create a UDP socket
2. **Assign a port to socket**
3. Communicate
4. Close the socket

## Server

1. Create a UDP socket
2. **Assign a port to socket**
3. Repeatedly
  - Communicate

# Example -Echo using Datagram socket

40

```
echoServAddr.sin_family = AF_INET; /* Internet address family */
echoServAddr.sin_addr.s_addr = inet_addr(echoServIP); /* Server IP address*/
echoServAddr.sin_port = htons(echoServPort); /* Server port */

echoStringLen = strlen(echoString); /* Determine input length */

/* Send the string to the server */
if (sendto(clientSock, echoString, echoStringLen, 0,
 (struct sockaddr *) &echoServAddr, sizeof(echoServAddr))
 != echoStringLen)
 DieWithError("send() sent a different number of bytes than expected");
```

## Client

1. Create a UDP socket
2. Assign a port to socket
3. **Communicate**
4. Close the socket

## Server

1. Create a UDP socket
2. **Assign a port to socket**
3. Repeatedly
  - Communicate



# Example -Echo using Datagram socket

41

```
for (;;) /* Run forever */
{
 clientAddrLen = sizeof(echoClientAddr) /* Set the size of the in-out parameter */
 /*Block until receive message from client*/
 if ((recvMsgSize = recvfrom(servSock, echoBuffer, ECHOMAX, 0),
 (struct sockaddr *) &echoClientAddr, sizeof(echoClientAddr))) < 0)
 DieWithError("recvfrom() failed");

 if (sendto(servSock, echobuffer, recvMsgSize, 0,
 (struct sockaddr *) &echoClientAddr, sizeof(echoClientAddr))
 != recvMsgSize)
 DieWithError("send() failed");
}
```

## Client

1. Create a UDP socket
2. Assign a port to socket
3. **Communicate**
4. Close the socket

## Server

1. Create a UDP socket
2. Assign a port to socket
3. Repeatedly
  - **Communicate**

# Example -Echo using Datagram socket

42

Similarly, the client receives the data from the server

## Client

1. Create a UDP socket
2. Assign a port to socket
3. **Communicate**
4. Close the socket

## Server

1. Create a UDP socket
2. Assign a port to socket
3. Repeatedly
  - **Communicate**

# Example -Echo using Datagram socket

43

```
close(clientSock);
```

## Client

1. Create a UDP socket
2. Assign a port to socket
3. Communicate
4. **Close the socket**

## Server

1. Create a UDP socket
2. Assign a port to socket
3. **Repeatedly**
  - Communicate